

Figure 1: Trajectories as recorded by GPS without any processing applied. All 15 datasets are shown.

## 1 Basic examination and processing

Within this strand of work we are looking at the GPS and IMU data from the RC car “long trajectories” in Stanmer Park. The overall goal is to examine this data, clean up any issues in it and make it ready for working on assessing visual navigation algorithms.

First, we examine the data by plotting the x and y coordinates of GPS locations taken along the 15 trajectories. The data is plotted in figures 1 to 3.

The first basic observations are:

1. It looks like the coordinates are in m - not in mm as stated in the csv files. If true, the trajectories are some 700m or so long. I will describe everything below under the assumption of units being m.
2. The trajectories start with the vehicle standing still but we can see some drift. Some trajectories have more, some less but on the order of tens of cm (see figure 2).
3. The trajectories presumably started at the same physical location but the recorded starting positions appear offset between trajectories by on the order of 4 – 8 m. Furthermore, normalising them to start from the same location does visually not seem to improve the overlap of trajectories, suggesting it is not just a global offset between them (see more about this below)
4. The end of the trajectories is fairly inconsistent, see also figure 3. To do meaningful comparisons of visual navigation algorithms between trajectories, it appears sensible to cut the trajectories to a common, fairly consistent part.

Based on the observations I have made the following first processing steps (`zero_chop.py`):




Figure 2: Detail of previous figure focussing on the trajectory start.

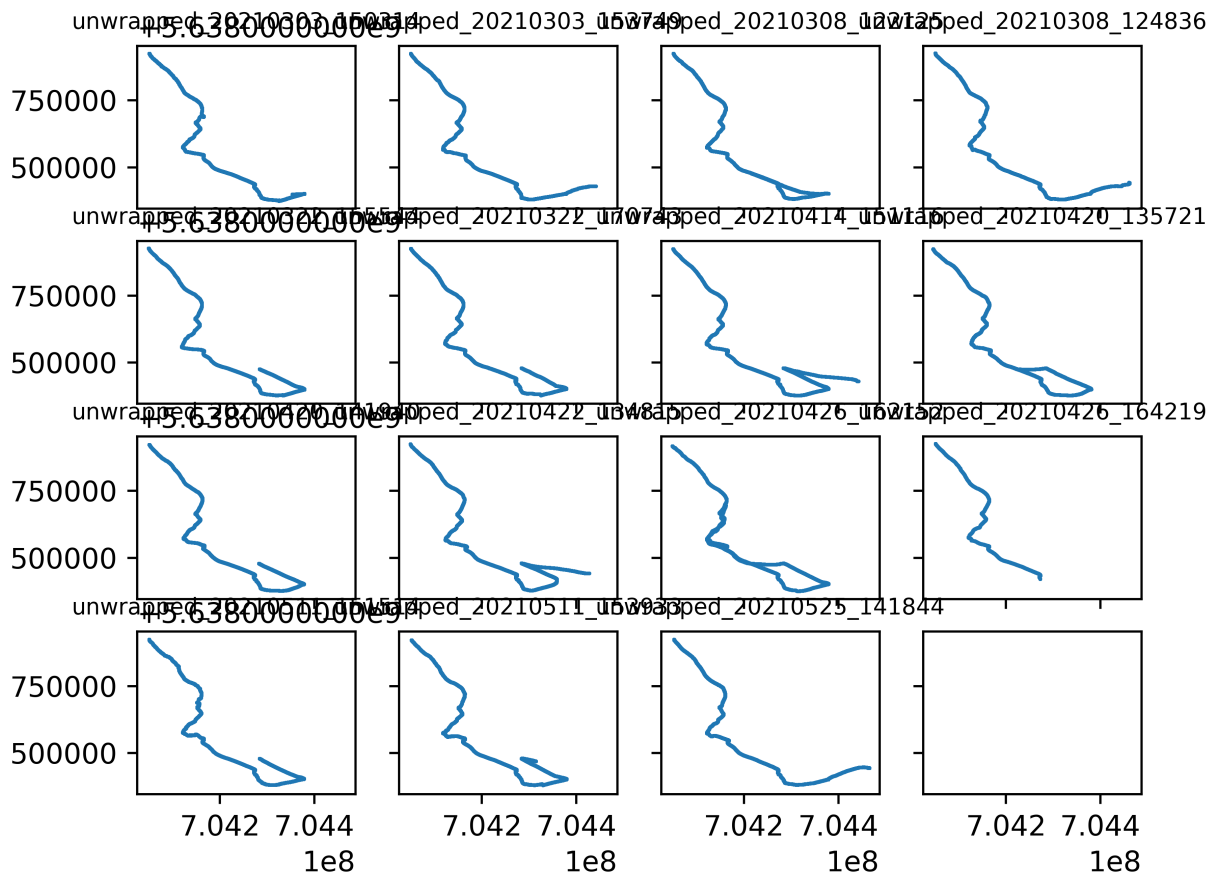


Figure 3: Trajectories as recorded by GPS without any processing applied. All 15 datasets are shown. Mainly included here to be able to identify trajectories and the corresponding data set names.




Figure 4: Minimally processed trajectories with initial non-moving phase cut and everything beyond the stated  $y$  threshold cut.

1. Remove the initial non-moving phase: Here I removed all points where the car did not move more than `min_move` m within `pt_dist` data points (or the next non-NaN point if that is a NaN point). For now I have used `min_move= 0.5` m and `pt_dist= 20`.
2. Remove any part of the trajectory once the  $y$  coordinate has passed below 5638423 for the first time. This cut-off was chosen by eyeballing the common part of all trajectories.

The resulting trajectories are shown in figure 4.

Before doing any more analysis, it is very inconvenient that about every third GPS position is recorded as NaN. This is presumably due to some mismatch of recording frequencies for GPS and other data and other GPS failures (there are also larger gaps at (thankfully) rare times). Also, there is some obvious noise in the GPS readings. To get something more consistent I have taken an approach of fitting pieces of the trajectory with a polynomial function and recombining all fitted pieces by averaging to one interpolated trajectory. This is done with the script `fit_trajectories.py`. For now, I have included results with polynomial degree 1 (simple linear regressions) and on pieces of trajectory of 75 data points each.

Next, we examine the heading information from fitted trajectories and from the IMU recordings. There are jumps in the IMU recordings, which are some kind of artefact of unknown origin. I was not able to find a simple explanation (such as power outage, with reset to 0 or such). I have removed these jumps by the following method (`correct_heading.py`):

1. Whenever the heading reading of the IMU jumps by more than 15 degrees in a single timestep, the jump is subtracted from the subsequent headings.
2. I shift the heading data globally so that the average IMU heading matches the average GPS-based heading.




Figure 5: Same a previous figure but with the initial position of the initial trajectory subtracted. Note how these do *not* appear to overlap better.




Figure 6: Heading shown as a function of time. For the purposes of this plot I have “unrolled” the angle so that there are no “jumps” of 360 degree.

GPS-based heading was calculated simply based on straight line connections between subsequent data points on the interpolated GPS trajectory.

Figure 6 shows the results. The original IMU data, also shifted to match means, is shown in blue, the corrected IMU data in orange and the GPS derived heading in green. The match is arguably sensible even though not perfect. For one, the GPS heading seems to undershoot in sharp turns. This could potentially be improved by taking less points for the trajectory regressions or a higher degree of the fitted polynomials. There are also a couple of sections in individual trajectories where the match is slightly less good.

To give a better impression of the quality of the fitted trajectory and heading fit, I have made movies that move a “magnifying glass” along the trajectories and display the original GPS locations, the fitted trajectory points and GPS derived and IMU derived headings (as a quiver plot). (`make_movie_frames.py` and then using `ffmpeg` like so: `>> ffmpeg -framerate 10 -pattern_type glob -i "frame_*.png" output.mp4`),

Figures 7 and 8 show two individual example frames of the movie. In the former, the agreement of the headings is excellent which is the case in most locations along the trajectories. In the second example, there is a bit more disagreement. However, this is much more rare and this is one of the more extreme examples. It is worth noting that the IMU headings after jump removal do not appear to degrade gradually along the trajectory but appear to be quite reliable to the very end. This suggests to me that they can be trusted to provide ground truth along the trajectory (and more so than the GPS-derived headings if there happens to be a discrepancy). Please note that I have here displayed the heading arrows next to the original GPS waypoints rather than the extrapolated points; this can be slightly confusing where it is not obvious which extrapolated point corresponds to which original waypoint.

Figure 9 shows the interpolated trajectories altogether. The original shape and spread of the trajectories is well-preserved (compare 4).

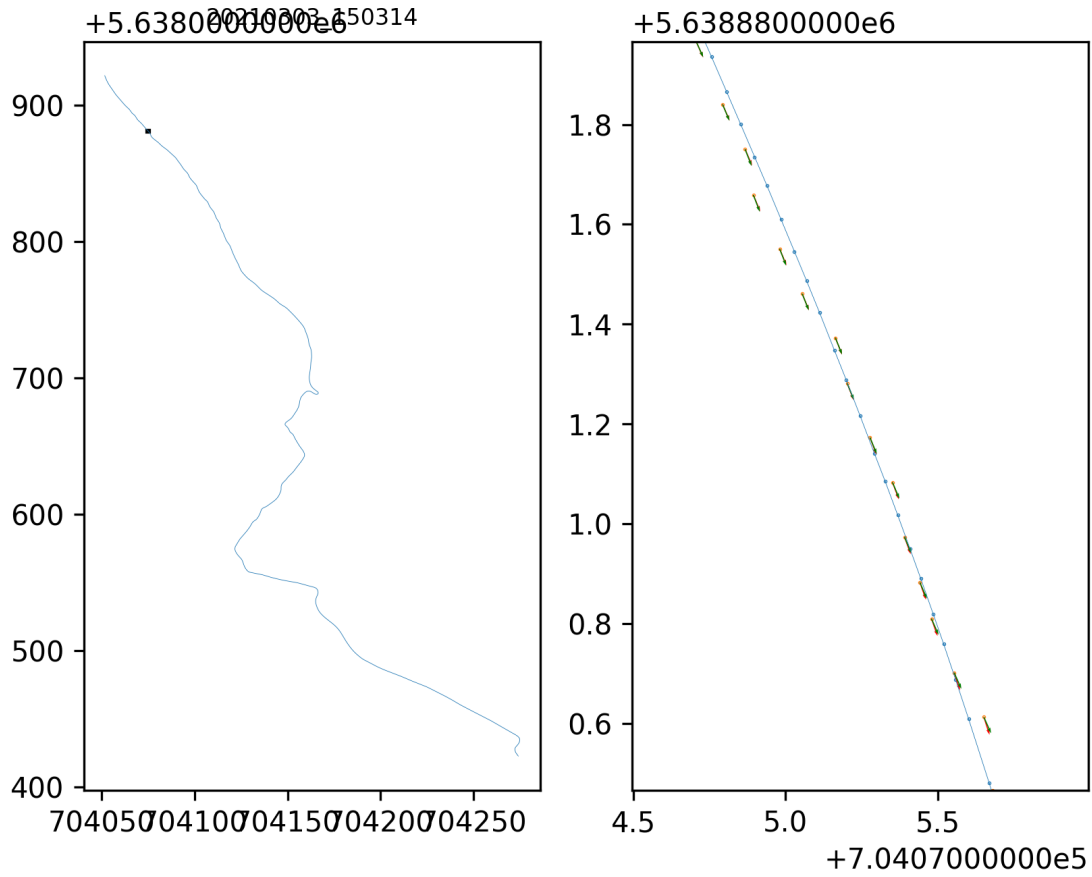


Figure 7: Example of the cleaned data. left: full trajectory, right: enlargement of the area marked by the small rectangle in the left panel. Orange dots are original GPS positions, blue dots and line the interpolated trajectory, and red arrows GPS heading as well as green arrows showing the IMU heading. In this example they coincide very well (which is true most of the time).

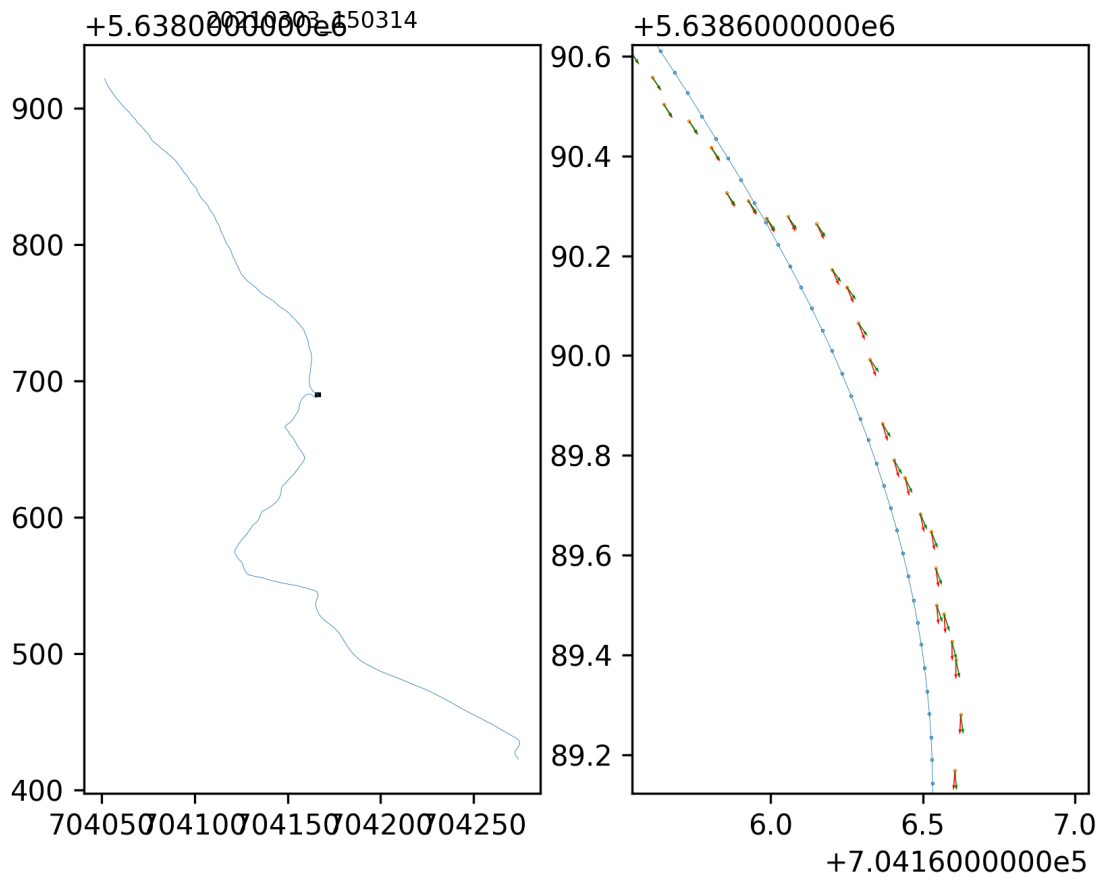


Figure 8: Example of the cleaned data as above. In this example the headings diverge more (this is one of the most extreme examples on this trajectory).






Figure 9: Fitted trajectories plotted together. The spread of the original data is reproduced quite accurately.




Figure 10: Same as previous figure but with starting points collapsed onto the same point  $(0, 0)$ .

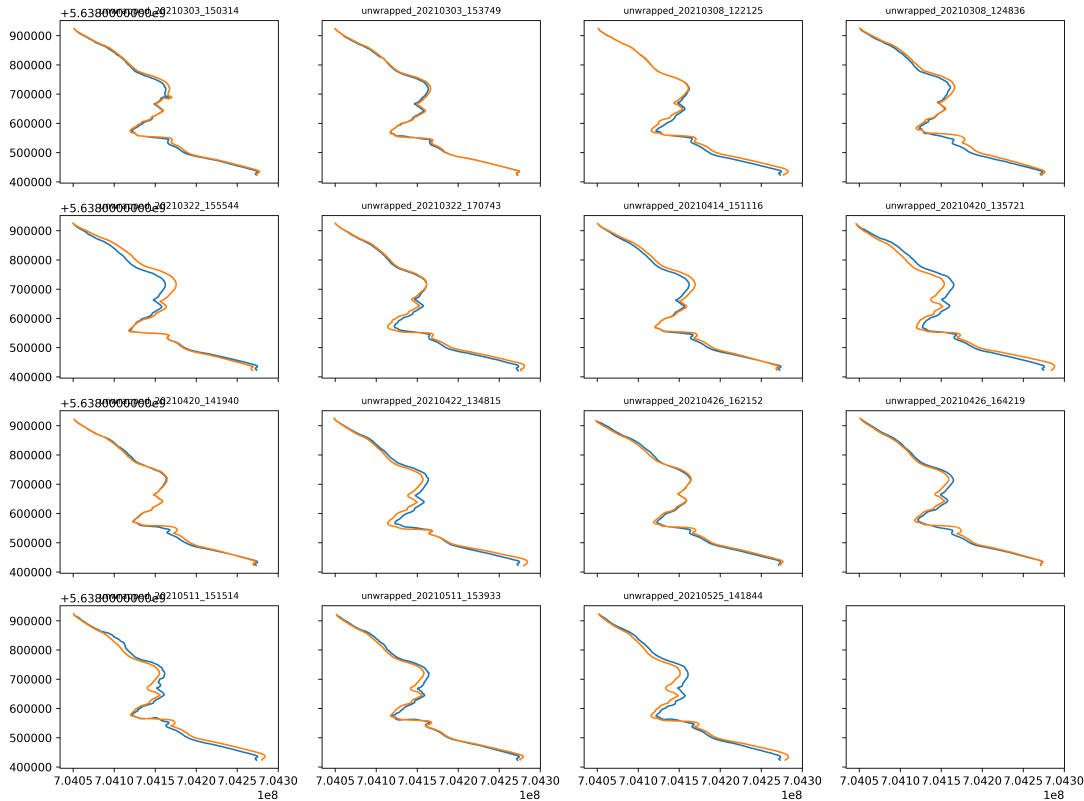


Figure 11: Trajectories reconstructed from IMU heading data (orange) compared to the interpolated trajectories (blue).

When I went ahead and measured the pairwise distances of interpolated trajectories with `measure_0.py`, the trajectories “as is” were closer together than if one makes the starting locations match (by subtracting them), compare figure 9 and 10. On average the summed distances were 26194.83 (raw) and 39009.99 (same starting point).

As a fun exercise I used the IMU headings to reconstruct trajectories. To lend a bit of a helping hand I used the interpolated  $x$  and  $y$  coordinates to determine the size of steps to make (the speed in other words). Use `reconstruct.py`. Figure 11 shows the results.

Even though there are obviously some discrepancies, overall the trajectories are very similar, confirming once again that the IMU heading data is quite accurate. There is a little bit of cheating involved here through providing “perfect speed information” but that does not take away from the fact that the heading data is overall mostly right.

## 2 Examining images

When observing odd trajectory bits in the generated trajectory movies, it is unclear what might have happened. We, therefore, visually inspected the camera footage of all datasets and identified obvious abnormalities. These were essentially, the robot being still because of a minor repair, or backing up briefly because stuck on something. We manually identified the following segments:

1. `unwrapped_20210303_150314` starts at 131; ends at tree 11304, no anomalies.
2. `unwrapped_20210303_153749` start at 125; ends at tree at about 10379 frames, no anomalies.

3. `unwrapped_20210308_122125` starts at 151; ends at tree at about 9700 frames, no anomalies.
4. `unwrapped_20210308_124836` starts at 148; in between it looks like a loose cable or a thin branch stuck on the robot; at 4209 strange maneuver to the left, back out, continue, normal again from 4338; 5336 driving against a tree, back out, continues route approx 5367; 8885 to 8906 slight stuck and back out maneuver; 9044 to 9074 stuck and back up; ends at tree about 10333.
5. `unwrapped_20210322_155544` starts at 383; 5420: taking a left turn because there are people in the way where he normally drives on the right of the prominent tree; reaches the tree at 9851 but then drives off into another direction without stopping
6. `unwrapped_20210322_170743` starts a 117; robot stuck at 7479, continues at 7608; reaches tree 9310, then drives off
7. `unwrapped_20210414_151116` start at 127; reaches tree 9157, then drives off
8. `unwrapped_20210420_135721` start at 339; 6294 to 6326 robot goes back and forth; same 6759 to 6781; 7138 to 7163; 7174 to 7209; 7834 to 7853; 8256 to 8280; 8371 to 8389; reaches tree at 10026 but the last bit of the route looked visually as if it was too far to the left!
9. `unwrapped_20210420_141940` start at 86; reaches tree 9410 then takes off.
10. `unwrapped_20210422_134815` start at 135; doesn't go to the tree! Instead to a different stopping point where it stands still for quite a while ...
11. `unwrapped_20210426_162152` starts at 50; reaches tree 3673 - very fast driving on this one!
12. `unwrapped_20210426_164219` starts at 69; frames stop prematurely at 3129; this is in agreement with the csv files.
13. `unwrapped_20210511_151514` starts at 199; 1377 to 1414 back and forth; from about frame 3635 to 3925 the robot stood still while Norbert was fixing the robot (apparently something had fallen off); reaches tree 10357
14. `unwrapped_20210511_153933` starts at 112; reaches tree at 10099, then drives off
15. `unwrapped_20210525_141844` starts at 154; stops at tree 8764 rest is irrelevant.

The manual exclusions were coded into a new script `zero_chop_visual.py` which operates on the data tables in the `unwrapped_*` directories. There are `fit_trajectories_visual.py` and `correct_heading_visual.py` scripts that essentially recapitulate the workflow above for the manually cut data. The results are essentially of similar quality, with some of the obvious artifacts removed.

The `make_movie_frames_visual.py` script improves on the previous script by including the camera feed into the movie so that one can now see the reconstructed trajectories and headings in conjunction with the views from the camera.

There is also a `reconstruct_visual.py` and `reconstruct_visual_greedy.py` script. The former recapitulates `reconstruct.py` but on the manually trimmed data sets. The `_greedy` version is trying to reconstruct the trajectory based on IMU data but rather than using the distance travelled between corresponding points on the fitted trajectory, it is determining the distance travelled that will minimise the distance of the IMU trajectory point to the corresponding fitted trajectory point. I named it greedy, as it is optimising this in a greedy way from the start of the trajectory and this local minimisation of divergence between trajectories can of course lead to larger divergences later on.

I have also tried a global optimisation using the Nelder-Mead simplex algorithm on the vector of all distances travelled but it appears that this problem might be too high-dimensional/ the optimisation to be too slow to be useful.

### 3 Appendix

To make the movies I used

```
ffmpeg -framerate 10 -pattern_type glob -i "frame_*.png" output.mp4
```